

Basic background



License & Disclaimer

2

License Information

This presentation is licensed under the
Creative Commons BY-NC License



To view a copy of the license, visit:

<http://creativecommons.org/licenses/by-nc/3.0/legalcode>

Disclaimer

- We disclaim any warranties or representations as to the accuracy or completeness of this material.
- Materials are provided “as is” without warranty of any kind, either express or implied, including without limitation, warranties of merchantability, fitness for a particular purpose, and non-infringement.
- Under no circumstances shall we be liable for any loss, damage, liability or expense incurred or suffered which is claimed to have resulted from use of this material.

Goal

3

- This module provides an overview of the basic background used in the Software Security Area. The module is structured in two lectures.
- In this first lecture we will learn...
 - how binary files are stored;
 - the basic steps that transform a source code into programs;
 - principles of x86 architecture;
 - mechanisms used to support memory management.

Prerequisites

4

- Lecture:
 - Basic knowledge of C
 - Basic knowledge of Python

Outline

5

- From Code to Programs
- Executable and Linkable Format (ELF)
- x86 architectures
- Memory management and Calling conventions
- Debugging

Outline

6

- From Code to Programs
- Executable and Linkable Format (ELF)
- x86 architectures
- Memory management and Calling conventions
- Debugging

From code to programs

7

- Compiling a C program is a multi-stage process composed of four steps:
 - preprocessing
 - compilation
 - assembly
 - linking

GCC Compiler

8

- In this lecture *gcc* compiler will be used...
 - is the C compiler included in the *GNU Compiler Collection*
 - is available in the main operating systems
 - Standard compiler in all the UNIX/Linux distributions
- Detailed documentation is available at the following link:

<https://gcc.gnu.org/onlinedocs/>

From code to programs: preprocessing

9

- In the first phase *preprocessor* commands (in C they start with '#') are interpreted:

```
#include <stdio.h>

#define MESSAGE "Hello world!"

int main() {
    printf(MESSAGE);
    return 0;
}
```

gcc -E hello.c

```
# 2 "hello.c" 2

# 5 "hello.c"
int main() {
    printf("Hello world!");
    return 0;
}
```

From code to programs: compilation

10

- In the second phase, preprocessed code is translated into *assembly instructions*:

```
#include <stdio.h>

#define MESSAGE "Hello world!"

int main() {
    printf(MESSAGE);
    return 0;
}
```

gcc -s hello.c

```
# 2 "hello.c" 2

# 5 "hello.c"
int main() {
    printf("Hello world!");
    return 0;
}
```

```
main:
.LFB0:
    .cfi_startproc
    pushq   %rbp
    .cfi_def_cfa_offset 16
    .cfi_offset 6, -16
    movq    %rsp, %rbp
    .cfi_def_cfa_register 6
    movl    $.LC0, %edi
    movl    $0, %eax
    call    printf
    movl    $0, %eax
    popq    %rbp
    .cfi_def_cfa 7, 8
    ret
```

From code to programs: assembly

11

- In the *assembly* phase assembly instructions are translated into *machine* or *object code*:

```
#include <stdio.h>

#define MESSAGE "Hello world!"

int main() {
    printf(MESSAGE);
    return 0;
}
```

gcc -c hello.c

```
# 2 "hello.c" 2

# 5 "hello.c"
int main() {
    printf("Hello world!");
    return 0;
}
```

```
main:
.LFB0:
.cfi_startproc
pushq %rbp
.cfi_def_cfa_offset 16
.cfi_offset 6, -16
movq %rsp, %rbp
.cfi_def_cfa_register 6
movl $.LC0, %edi
movl $0, %eax
call printf
movl $0, %eax
popq %rbp
.cfi_def_cfa 7, 8
ret
```



hello.o

From code to programs: linking

12

- In the last phase (multiple) *object code* are combined in a single executable
- In the generated file, references (links) to the used library are added



Static vs Dynamic Linking

13

- Two approaches can be used in the linking phase:
 - **Static Link**
 - Binaries are *self-contained* and do not depend on any external libraries
 - **Dynamic Link**
 - Binaries rely on system libraries that are loaded when needed
 - Mechanisms are needed to *dynamically* relocate code

Outline

14

- From Code to Programs
- Executable and Linkable Format (ELF)
- x86 architectures
- Memory management and Calling conventions
- Debugging

Executable and Linkable Format

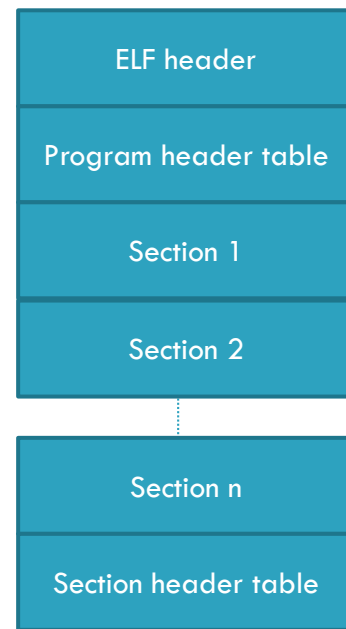
15

- The Executable and Linkable Format (ELF) is a common file format for object files
- There are three types of object files:
 - **Relocatable file** containing code and data that can be linked with other object files to create an **executable** or shared object file
 - **Executable files** holding a program suitable for execution
 - **Shared object files** that can be:
 - linked with other relocatable and shared object files to obtain another object file
 - used by a **dynamic linker** together with other executable files and object files to create a **process image**

Executable and Linkable Format

16

- Any ELF file is structured as:
 - an **ELF header** describing the file content
 - a **Program header table** providing info about how to create a process image
 - a sequence of **Sections** containing what is needed for linking (instructions, data, symbol table, relocation information,...)
 - a **Section header table** with a description of previous sections



ELF: Relevant sections

17

- **.text**: contains the executable instructions of a program
- **.bss**: contains uninitialised data that contribute to the program's memory image
- **.data, .data1**: contain initialized data that contribute to the program's memory image
- **.rodata, .rodata1**: are similar to **.data** and **.data1**, but refer to read-only data
- **.symtab**: contains the program's symbol table
- **.dynamic**: provides linking information

Outline

18

- From Code to Programs
- Executable and Linkable Format (ELF)
- **x86 architectures**
- Memory management and Calling conventions
- Debugging

x86 Instruction Sets

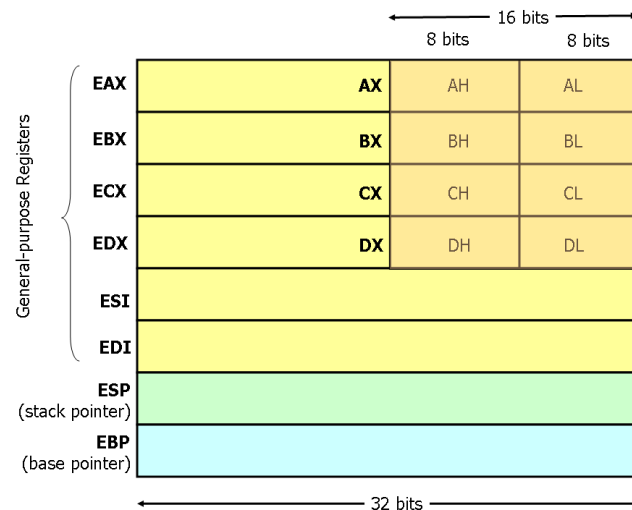
19

- We provide a short introduction of a small but useful subset of the available instructions and assembler directives of x86 assembly language
- A detailed description can be found at the following links:
 - Intel x86 Instruction Set Reference
 - <http://www.felixcloutier.com/x86/>
 - Intel's Pentium Manuals
 - <http://www.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html>

X86-32 Registers

20


- x86-32 processors have eight 32-bit general purpose registers
- The register names are mostly historical...
 - *EAX* used to be called *the accumulator* since it was used by a number of arithmetic operations
 - *ECX* was known as the *counter* since it was used to hold a loop index
- Two are reserved for special purposes:
 - the *stack pointer* (ESP)
 - the *base pointer* (EBP)



X86-32 Registers

21

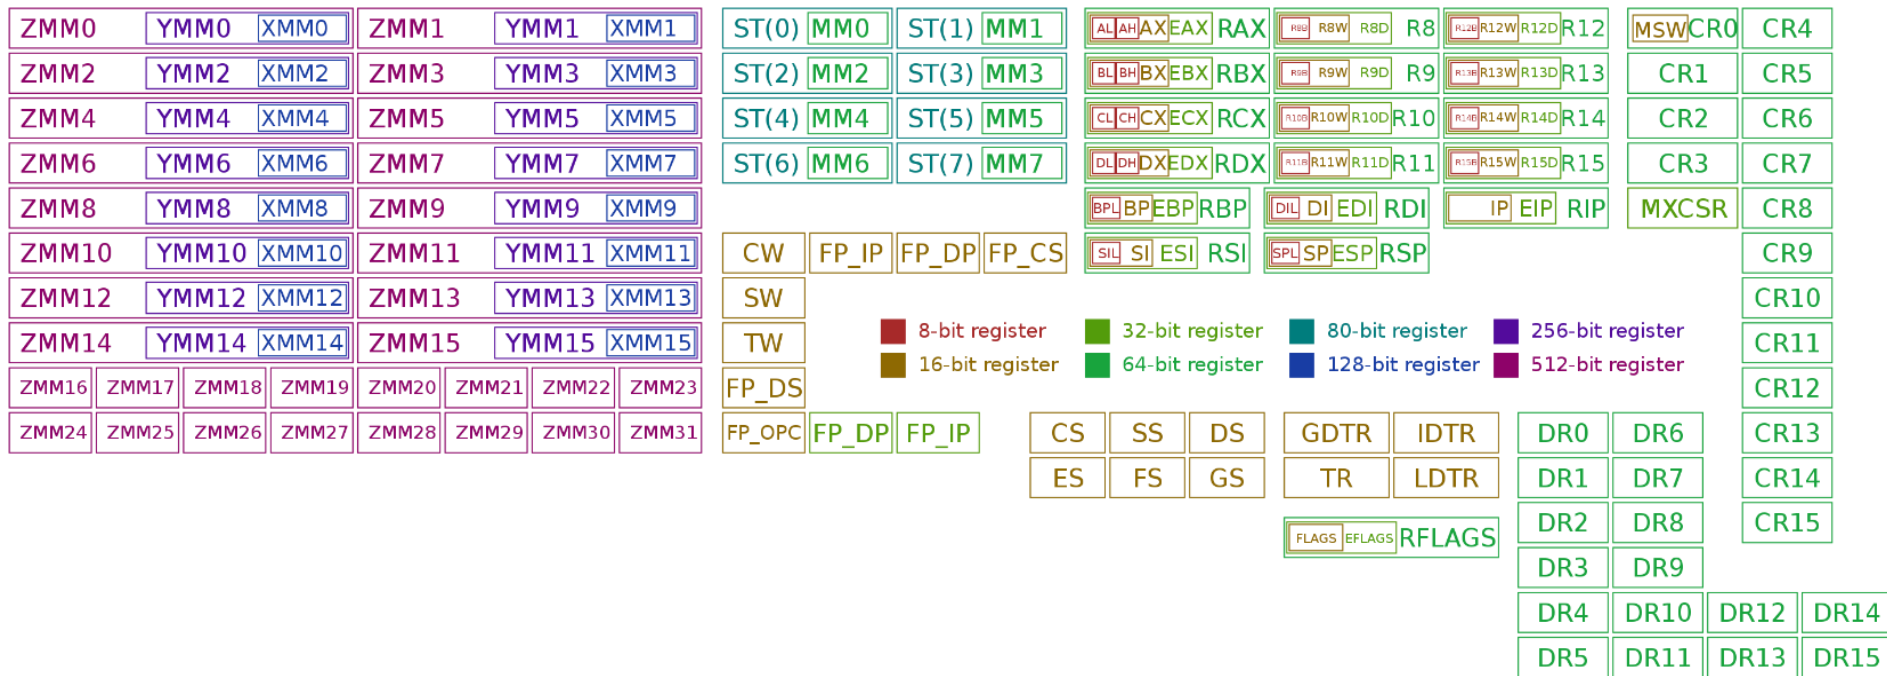
- For the *EAX*, *EBX*, *ECX*, and *EDX* registers, subsections may be used
- For example:
 - *AX* refers to the least significant 2 bytes of *EAX*
 - *AL* refers to the least significant byte of *AX*
 - *AH* refers to the most significant byte of *AX*
- These sub-registers are mainly hold-overs from older, 16-bit versions of the instruction set



X86-64 Registers

22

- X86-64 architecture provides a larger set of registers:



x86 Instructions

23

- Data Movement Instructions
 - *mov op1,op1*: copies the data item referred to by *op1* into the location referred to by *op2*
 - *push op1*: places its operand onto the top of the stack
 - *pop op1*: removes the 4-byte data element from the top of the stack
 - *lea op1,op2*: load the memory address indicated by *op2* into the register specified by *op1*

x86 Instructions

24

➤ Arithmetic and Logic Instructions

➤ *add op1,op2*: stores in *op1* the result of *op2+op1*

➤ *sub op1,op2*: stores in *op1* the result of *op2-op1*

➤ *and op1,op2*

➤ *or op1,op2*

➤ *xor op1,op2*

➤ ...

} Perform the specified logical operation on the operands, storing the result in the first operand location

x86 Instructions

25

- Control Flow Instructions
 - *jmp op*: jump to the instruction at the memory location specified by the operand *op*
 - *cmp op1,op2*: compares the values of the two specified operands and stores the result in the *machine status word*
 - *j<condition> op*: depending on the *<condition>* and on the context of *machine status word*, jumps to instruction at the memory location indicated by the operand

Outline

26

- From Code to Programs
- Executable and Linkable Format (ELF)
- x86 architectures
- **Memory management and Calling conventions**
- Debugging

Memory management

27

- Many of the modern programming languages allow programmers to use **data types** without having to know **how** they are represented
- Similarly, programmers often ignore **where** data is stored (**allocated**)...
 - compiler makes decisions, however at runtime allocation is under the control of Operating System and CPU

Memory management

28

- In some programming languages, like C, memory management can be controlled by programmers:
 - memory can be **dynamically** allocated and deallocated
 - memory address of variables can be obtained (pointers)
- If x is a variable, $\&x$ denotes **the pointer to x** , i.e., the memory address where x is stored

Memory allocation...

29

- Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

We can assume that:

- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes

Memory allocation...

30

- Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;
    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

Variable declarations

We can assume that:

- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes

Memory allocation...

31

- Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

Variable addresses


We can assume that:

- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes


Memory allocation...

32

- Memory is usually represented as a sequence of bytes:



In a $n \cdot 8$ architecture, bytes are arranged in groups of n :



Memory allocation...

33

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

We can run this simple program to observe how data are allocated in memory

Memory allocation...

34

- Different allocation policies can be used by *gcc*:

```
[CC> gcc -o alignment alignment.c  
[CC> ./alignment  
i is allocated at 0x7fffee117e7cc  
c is allocated at 0x7fffee117e7cb  
s is allocated at 0x7fffee117e7c8  
l is allocated at 0x7fffee117e7c0  
[CC>
```

```
[CC> gcc -o alignment alignment.c -O2  
[CC> ./alignment  
i is allocated at 0x7fffee4dbb7c8  
c is allocated at 0x7fffee4dbb7cf  
s is allocated at 0x7fffee4dbb7cc  
l is allocated at 0x7fffee4dbb7c0  
[CC>
```

Data alignment


35

- Compilers may introduce **padding** or change the order of data in memory
- There are trade-offs between speed and memory usage
- C compilers provide many optional optimization

Memory segments

36

- Memory is allocated for each **process** (a running program) to store **data** and **code**.
- This allocated memory consists of different **segments**:
 - **stack**: for local variables
 - **heap**: for dynamic memory
 - **data segment**:
 - *global uninitialized variables (.bss)*
 - *global initialized variables (.data)*
 - **code segment**



The stack

37

- The stack consists of a sequence of *stack frames* (or activation records), each for each function call:
 - allocated on *call*
 - de-allocated on *return*

```
int main(int argc, char **argv) {  
    int f = fib( n: 10);  
    printf("FIB(10)=%d\n",f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib( n: n-1);  
        f2 = fib( n: n-2);  
        return f1+f2;  
    }  
}
```

stack frame
for main()

stack frame
for fib()

stack frame
for fib()

Unused memory

The stack

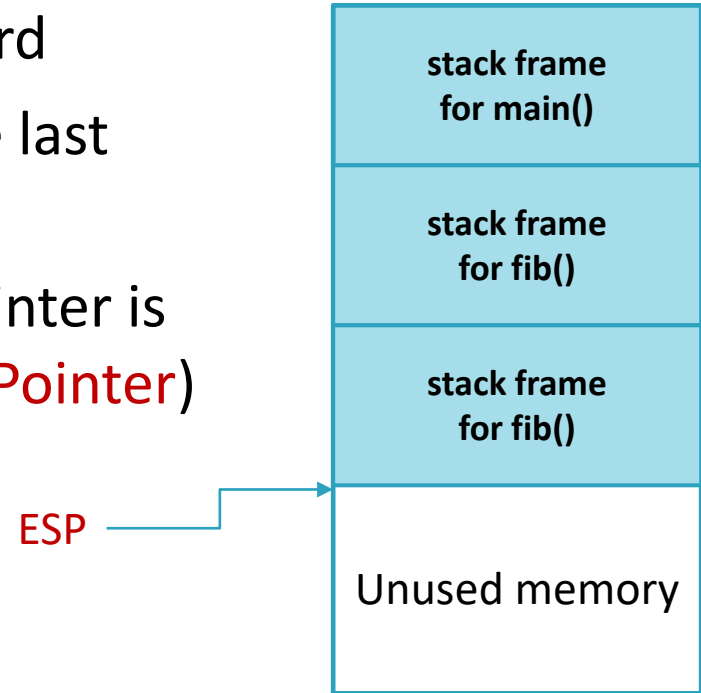
38

- The precise structure and organization of the stack depends on system architecture, operating system, and compilers we are using.
- For the sake of simplicity, in this lecture we will focus on x86 architectures (32 and 64 bits) and on *gcc* compiler on *Linux*
- More detailed information are available at the following links:
 - <https://docs.microsoft.com/en-us/cpp/cpp/argument-passing-and-naming-conventions>
 - <http://refspecs.linuxbase.org/>

The stack

39

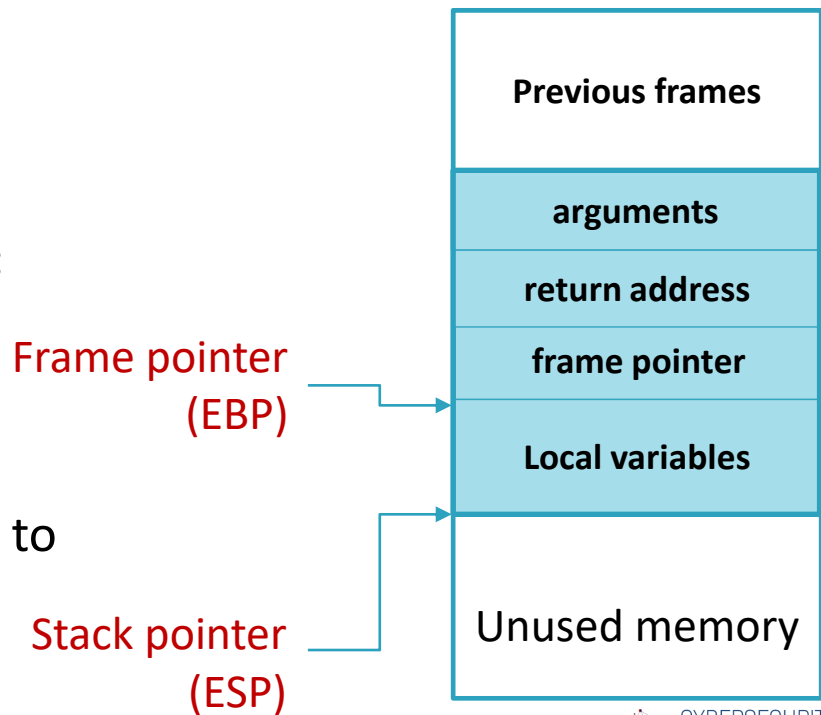
- Typically, the stack grows downward
- The **stack pointer (SP)** refers to the last element on the stack
- On x86 architectures, the stack pointer is stored in the **ESP (Extended Stack Pointer)** register



Stack frame (for x86)

40

- In x86 architecture, each stack frame contains:
 - Function arguments
 - Local variables
 - Copies of registries that must be restored:
 - return address
 - previous frame pointer
- Frame pointer, named Extended Base Pointer (EBP), provides a starting point to local variables




Stack frame: example

41

```
int main(int argc, char **argv) {  
    int f = fib( n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib( n: n-1);  
        f2 = fib( n: n-2);  
        return f1+f2;  
    }  
}
```

Frame pointer
Stack pointer

Function fib is
invoked with
parameter 10



Stack frame: example


42

```
int main(int argc, char **argv) {  
    int f = fib( n:10);  
    printf("FIB(10)=%d\n",f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib( n: n-1);  
        f2 = fib( n: n-2);  
        return f1+f2;  
    }  
}
```

Stack frame is
allocated and
pointers updated

Frame pointer

Stack pointer




Stack frame: example

43

```
int main(int argc, char **argv) {  
    int f = fib( n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib( n: n-1);  
        f2 = fib( n: n-2);  
        return f1+f2;  
    }  
}
```

Frame pointer →
Stack pointer →

When a function returns, pointers are updated. Function result (if any) is copied in a register



Stack frame (for x86-64)

44

- In the *64-bit* version of x86:
 - Registers are extended to 64 bits
 - A new set of 8 registers is added
- The arguments are not all in the stack:
 - The first 6 are passed via registers
 - The remaining are placed in the stack (like for x86)
- Pointers *EBP* and *ESP* are named *RBP* and *RSP*
- A *red zone* of 128 bytes is placed in the stack just under RSP
 - This can be used to store extra local variables and is not modified by interrupt/exception/signal handlers

C declaration: cdecl

45

- The **cdecl** (C declaration) is a **calling convention** used in many compilers for x86
 - A calling convention describes how functions receive their parameters from caller and how they return their results
- In cdecl arguments are passed on the stack, while values are returned via registers
- Function arguments are pushed on the stack in the *right-to-left order*
- The caller *cleans* the stack after the RETURN from the called function

The heap

46

- Memory allocation and de-allocation in the stack is very fast
 - However, this memory cannot be used after a function returns
- The heap is used to store dynamically allocated data that outlive function calls:
 - This area is under programmer's responsibility

Memory management functions

47

- Basic C functions for memory management are:
 - *malloc(int)*, given an integer n allocates an area of n (continuous) bytes and returns a **pointer** to that area
 - *free(void*)*, deallocates the memory associated with a pointer

Stack and Heap pointers

48

- We have always to preserve reference to allocated areas in the heap
 - Otherwise, we cannot use them!
- These references (pointers) can be stored in the stack or in the heap
- We could store pointers to the stack in the heap
 - This can be dangerous! Referenced memory could be released!

Outline

49

- From Code to Programs
- Executable and Linkable Format (ELF)
- x86 architectures
- Memory management and Calling conventions
- **Debugging**

Debugging

50

- Debugging is the process of finding and fixing bugs (defects or problems that prevent correct operation) within computer programs, software, or systems.
- Debugging can be based on different techniques and methodologies such as:
 - interactive debugging
 - control flow analysis
 - unit and integration testing
 - log file and output analysis
 - memory dumps
 - profiling
- Many programming languages and software development tools also offer programs to aid in debugging, known as *debuggers*

Debugging

51

- A **debugger** is a software tool that allows to:
 - run the target program under controlled conditions
 - track program operations in progress
 - monitor changes in computer resources
 - display the contents of memory
 - modify memory or register contents

Debugging

52

- Compilers can be instrumented to emit extra (optional) data that debugger can use for a more informative input
- In the case of *gcc* compiler, the parameter *-g* can be used
- Debugging information is stored in specific sections of ELF file:
 - *.debug*, containing info for symbolic debugging
 - *.line*, containing line number information

Example

53

```
CC> gcc -o hello -g hello.c
CC> readelf --debug-dump hello
Contents of the .debug_aranges section:
```

```
Length:                44
Version:                2
Offset into .debug_info: 0x0
Pointer Size:           8
Segment Size:           0
```

Address	Length
0000000000400526	000000000000001a
0000000000000000	0000000000000000

```
Contents of the .debug_info section:
```

```
Compilation Unit @ offset 0x0:
Length:          0x8d (32-bit)
Version:         4
Abbrev Offset:   0x0
```

Basic background

